


## MySQL Exercise 8: Joining Tables with Outer Joins

We focused on inner joins in the last set of exercises. Most of the time inner joins will give you the results you are looking for. Occasionally, though, you might want to include all the data from a table in your calculations or your output, even if those data do not all match up with the data from the other tables you are joining with.

For example, if you have a table with customer demographic information and a table with information about which customers were sent a free sample, you might want to analyze the characteristics of customers who did and did not receive the sample. The best way to do that in a program like Tableau would be to have all your customer information in one table with an extra column indicating whether the customer received the free sample or not. Alternatively, you might want to generate a list of customers who did *not* receive the free sample, so that you can arrange for the customers to receive a free sample in the future.

In these types of situations, you will use outer joins to connect tables. Outer joins include left joins, right joins, or full outer joins (recall that full outer joins are NOT supported in MySQL). Refer to the videos "What are Joins?" and "Joins With Many to Many Relationships and Duplicates" for more information about joins.

Here's a picture to remind you of the general concepts behind outer joins:



To begin practicing outer joins, load the sql library, connect to the Dognition database, and make the Dognition database your default database:

```
In [4]:
%load_ext sql
%sql mysql://studentuser:studentpw@localhost/dognitiondb
%sql USE dognitiondb

* mysql://studentuser:***@localhost/dognitiondb
0 rows affected.

Out[4]:
[]
```

## Left and Right Joins

Left and right joins use a different syntax than we used in the lesson about inner joins. The method I showed you to execute inner joins tells the database how to relate tables in a WHERE clause like this:

```
WHERE d.dog_guid=r.dog_guid
```

I find this syntax -- called the "equijoin" syntax -- to be very intuitive, so I thought it would be a good idea to start with it. However, we can re-write the inner joins in the same syntax used by outer joins. To use this more traditional syntax, you have to tell the database how to connect the tables using an ON clause that comes right after the FROM clause. Make sure to specify the word "JOIN" explicitly. This traditional version of the syntax frees up the WHERE clause for other things you might want to include in your query. Here's what one of our queries from the inner join lesson would look like using the traditional syntax:

```
SELECT d.dog_guid AS DogID, d.user_guid AS UserID, AVG(r.rating) AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM dogs d JOIN reviews r
ON d.dog_guid=r.dog_guid AND d.user_guid=r.user_guid
GROUP BY d.user_guid
HAVING NumRatings > 9
ORDER BY AvgRating DESC
LIMIT 200
```

You could also write "INNER JOIN" instead of "JOIN" but the default in MySQL is that JOIN will mean inner join, so including the word "INNER" is optional.

If you need a WHERE clause in the query above, it would go after the ON clause and before the GROUP BY clause.

Here's an example of a different query we used in the last lesson that employed the equijoin syntax:

```
SELECT d.user_guid AS UserID, d.dog_guid AS DogID,
       d.breed, d.breed_type, d.breed_group
FROM dogs d, complete_tests c
WHERE d.dog_guid=c.dog_guid AND test_name='Yawn Warm-up';
```

**Question 1: How would you re-write this query using the traditional join syntax?**

In [5]:

```
%%sql
SELECT d.dog_guid AS UserID, d.user_guid AS UserID, d.breed, d.breed_type, d.breed_group
FROM dogs d JOIN complete_tests c
ON d.dog_guid=c.dog_guid
WHERE test_name="Yawn Warm-up"
LIMIT 10;
```

```
* mysql://studentuser:***@localhost/dognitiondb
10 rows affected.
```

Out[5]:

UserID	UserID_1	breed	breed_type	breed_group
fd27b86c-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	Pure Breed	Toy
fd27bbbe-7144-11e5-ba71-058fbc01cf0b	ce135f2c-7144-11e5-ba71-058fbc01cf0b	Mixed	Mixed Breed/ Other/ I Don't Know	None
fd27ba1a-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	Pure Breed	Toy
fd27c5be-7144-11e5-ba71-058fbc01cf0b	ce136ac6-7144-11e5-ba71-058fbc01cf0b	Shih Tzu-Poodle Mix	Cross Breed	None
fd27b948-7144-11e5-ba71-058fbc01cf0b	ce13615c-7144-11e5-ba71-058fbc01cf0b	Siberian Husky	Pure Breed	Working
fd27b5ba-7144-11e5-ba71-058fbc01cf0b	ce1353d8-7144-11e5-ba71-058fbc01cf0b	Shetland Sheepdog	Pure Breed	Herding
fd27c852-7144-11e5-ba71-058fbc01cf0b	ce136ee0-7144-11e5-ba71-058fbc01cf0b	Pug	Pure Breed	Toy
fd27c8d4-7144-11e5-ba71-058fbc01cf0b	ce136f94-7144-11e5-ba71-058fbc01cf0b	Boxer	Pure Breed	Working
fd27cf28-7144-11e5-ba71-058fbc01cf0b	ce13750c-7144-11e5-ba71-058fbc01cf0b	Chesapeake Bay Retriever	Pure Breed	Sporting
fd27c7d0-7144-11e5-ba71-058fbc01cf0b	ce136e36-7144-11e5-ba71-058fbc01cf0b	Vizsla	Pure Breed	Sporting

## CODE

```
%%sql SELECT d.dog_guid AS UserID, d.user_guid AS UserID, d.breed, d.breed_type, d.breed_group FROM dogs d JOIN complete_tests c ON d.dog_guid=c.dog_guid WHERE test_name="Yawn Warm-up";
```

## OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 20845 rows affected.

Now that we've seen the join syntax, we can begin practicing outer joins. Our Dognition data set will make outer joins more challenging than usual, due to the lack of declared primary keys in the original database, the many-to-many relationships, and the presence of duplicate rows and NULL values in columns we will be using to combine tables. Mastering outer joins in this challenging context, though, will ensure that you understand the fundamental concepts behind joins, which will be a terrific benefit when you start writing queries in other company databases.

The first query I described above was originally written to address the question of whether dog owners who are particularly surprised by their dog's performance on Dognition tests tend to own similar breeds of dogs. When we designed this query in the last lesson, we wanted to focus on the dog owners who reported the highest average amount of surprise at their dog's performance, and provided at least 10 ratings for one or more of their dogs in the ratings. We also wanted to examine the breed, breed\_type, and breed\_group of each of these owner's dogs.

In examining the query, we learned that:

- All of the user\_guids in the reviews table are in the dogs table
- Only 389 of the over 5000 dog\_guids in the reviews table are in the dogs table

The inner join we executed resulted in 389 rows of output, because it only included the data from rows that have equivalent values in both tables being joined. But what if we wanted the full list of dogs in the reviews table and an indication of whether or not they were in the dogs table, rather than only a list of review information from dogs in the dogs table? To achieve this list, we could execute an outer join.

Let's start by using a left outer join to get the list we want. When we use the traditional join syntax to write inner joins, the order you enter the tables in your query doesn't matter. In outer joins, however, the order matters a lot. A left outer join will include all of the rows of the table to the left of the LEFT JOIN keywords. A right outer join will include all of the rows of the table to the right of the RIGHT JOIN keywords. So in order to retrieve a full list of dogs who completed at least 10 tests in the reviews table, and include as much breed information as possible, we could query:

```
SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating) AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM reviews r LEFT JOIN dogs d
  ON r.dog_guid=d.dog_guid AND r.user_guid=d.user_guid
WHERE r.dog_guid IS NOT NULL
GROUP BY r.dog_guid
HAVING NumRatings >= 10
ORDER BY AvgRating DESC;
```

**Question 2: How could you retrieve this same information using a RIGHT JOIN?**

In [8]:

```
%%sql
SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating)
AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM reviews r LEFT JOIN dogs d
ON r.dog_guid=d.dog_guid AND r.user_guid=d.user_guid
WHERE r.dog_guid IS NOT NULL
GROUP BY r.dog_guid
HAVING NumRatings >=10
ORDER BY AvgRating DESC
LIMIT 10;
```

\* mysql://studentuser:\*\*\*@localhost/dognitiondb  
10 rows affected.

Out[8]:

rDogID	dDogID	rUserID	dUserID	AvgRating	NumRatings	breed	breed_group	breed_type
fdbf39f8-7144-11e5-ba71-058fbc01cf0b	fdbf39f8-7144-11e5-ba71-058fbc01cf0b	ce987914-7144-11e5-ba71-058fbc01cf0b	ce987914-7144-11e5-ba71-058fbc01cf0b	8.0000	12	Canaan Dog	Herding	Pure Breed
ce47553e-7144-11e5-ba71-058fbc01cf0b	None	ce6ca9ba-7144-11e5-ba71-058fbc01cf0b	None	7.8750	16	None	None	None
ce6f07e6-7144-11e5-ba71-058fbc01cf0b	None	ce7091e2-7144-11e5-ba71-058fbc01cf0b	None	7.5000	10	None	None	None
ce45ae5a-7144-11e5-ba71-058fbc01cf0b	None	ce67562c-7144-11e5-ba71-058fbc01cf0b	None	7.3529	17	None	None	None
ce2a68ac-7144-11e5-ba71-058fbc01cf0b	None	ce2a45c0-7144-11e5-ba71-058fbc01cf0b	None	7.1333	15	None	None	None
ce72be0e-7144-11e5-ba71-058fbc01cf0b	None	ce73f7a6-7144-11e5-ba71-058fbc01cf0b	None	7.0000	12	None	None	None
ce93d206-7144-11e5-ba71-058fbc01cf0b	None	ce8be19a-7144-11e5-ba71-058fbc01cf0b	None	6.7273	11	None	None	None
ce97cd7a-7144-11e5-ba71-058fbc01cf0b	None	ce948926-7144-11e5-ba71-058fbc01cf0b	None	6.7273	11	None	None	None
ce7038dc-7144-11e5-ba71-058fbc01cf0b	None	ce71a230-7144-11e5-ba71-058fbc01cf0b	None	6.6667	18	None	None	None
ce7dc3da-7144-11e5-ba71-058fbc01cf0b	None	ce7cbb2-7144-11e5-ba71-058fbc01cf0b	None	6.3750	16	None	None	None

CODE

```
%%sql SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating) AS AvgRating,
COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type FROM reviews r LEFT JOIN dogs d ON r.dog_guid=d.dog_guid AND
r.user_guid=d.user_guid WHERE r.dog_guid IS NOT NULL GROUP BY r.dog_guid HAVING NumRatings >=10 ORDER BY AvgRating DESC;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 932 rows affected.

QUESTION 1 ANSWER

In [14]:

```
%%sql
SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating)
AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM dogs d RIGHT JOIN reviews r
ON d.dog_guid=r.dog_guid AND d.user_guid=r.user_guid
WHERE r.dog_guid IS NOT NULL
GROUP BY r.dog_guid
HAVING NumRatings >=10
ORDER BY AvgRating DESC
LIMIT 10;
```

```
* mysql://studentuser:***@localhost/dognitiondb
10 rows affected.
```

Out[14]:

rDogID	dDogID	rUserID	dUserID	AvgRating	NumRatings	breed	breed_group	breed_type
fdbf39f8-7144-11e5-ba71-058fbc01cf0b	fdbf39f8-7144-11e5-ba71-058fbc01cf0b	ce987914-7144-11e5-ba71-058fbc01cf0b	ce987914-7144-11e5-ba71-058fbc01cf0b	8.0000	12	Canaan Dog	Herding	Pure Breed
ce47553e-7144-11e5-ba71-058fbc01cf0b	None	ce6ca9ba-7144-11e5-ba71-058fbc01cf0b	None	7.8750	16	None	None	None
ce6f07e6-7144-11e5-ba71-058fbc01cf0b	None	ce7091e2-7144-11e5-ba71-058fbc01cf0b	None	7.5000	10	None	None	None
ce45ae5a-7144-11e5-ba71-058fbc01cf0b	None	ce67562c-7144-11e5-ba71-058fbc01cf0b	None	7.3529	17	None	None	None
ce2a68ac-7144-11e5-ba71-058fbc01cf0b	None	ce2a45c0-7144-11e5-ba71-058fbc01cf0b	None	7.1333	15	None	None	None
ce72be0e-7144-11e5-ba71-058fbc01cf0b	None	ce73f7a6-7144-11e5-ba71-058fbc01cf0b	None	7.0000	12	None	None	None
ce93d206-7144-11e5-ba71-058fbc01cf0b	None	ce8be19a-7144-11e5-ba71-058fbc01cf0b	None	6.7273	11	None	None	None
ce97cd7a-7144-11e5-ba71-058fbc01cf0b	None	ce948926-7144-11e5-ba71-058fbc01cf0b	None	6.7273	11	None	None	None
ce7038dc-7144-11e5-ba71-058fbc01cf0b	None	ce71a230-7144-11e5-ba71-058fbc01cf0b	None	6.6667	18	None	None	None
ce7dc3da-7144-11e5-ba71-058fbc01cf0b	None	ce7cbb2a-7144-11e5-ba71-058fbc01cf0b	None	6.3750	16	None	None	None

CODE

```
%%sql SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating) AS AvgRating,
COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type FROM dogs d RIGHT JOIN reviews r ON d.dog_guid=r.dog_guid AND
d.user_guid=r.user_guid WHERE r.dog_guid IS NOT NULL GROUP BY r.dog_guid HAVING NumRatings >=10 ORDER BY AvgRating DESC;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 932 rows affected.

Notice in the output of both the left and the right version of the outer join, all the rows that had a dog\_guid in the reviews table but did NOT have a matching dog\_guid in the dogs table have the word "None" entered in output columns related to the dogs table. "None", in this case, is Jupyter's way of saying the value is NULL. This becomes clear when you query a list of only the dog\_guids that were NOT in the dogs table:

```
SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.r
ating) AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM reviews r LEFT JOIN dogs d
ON r.dog_guid=d.dog_guid AND r.user_guid=d.user_guid
WHERE d.dog_guid IS NULL
GROUP BY r.dog_guid
HAVING NumRatings >= 10
ORDER BY AvgRating DESC;
```

Go ahead and try it yourself (you should get 894 rows in your query):

In [18]:

```
%%sql
SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating)
AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM reviews r LEFT JOIN dogs d
ON r.dog_guid=d.dog_guid AND r.user_guid=d.user_guid
WHERE d.dog_guid IS NULL
GROUP BY r.dog_guid
HAVING NumRatings >=10
ORDER BY AvgRating DESC
LIMIT 10;
```

```
* mysql://studentuser:***@localhost/dognitiondb
10 rows affected.
```

Out[18]:


rDogID	dDogID	rUserID	dUserID	AvgRating	NumRatings	breed	breed_group	breed_type
ce47553e-7144-11e5-ba71-058fbc01cf0b	None	ce6ca9ba-7144-11e5-ba71-058fbc01cf0b	None	7.8750	16	None	None	None
ce6f07e6-7144-11e5-ba71-058fbc01cf0b	None	ce7091e2-7144-11e5-ba71-058fbc01cf0b	None	7.5000	10	None	None	None
ce45ae5a-7144-11e5-ba71-058fbc01cf0b	None	ce67562c-7144-11e5-ba71-058fbc01cf0b	None	7.3529	17	None	None	None
ce2a68ac-7144-11e5-ba71-058fbc01cf0b	None	ce2a45c0-7144-11e5-ba71-058fbc01cf0b	None	7.1333	15	None	None	None
ce72be0e-7144-11e5-ba71-058fbc01cf0b	None	ce73f7a6-7144-11e5-ba71-058fbc01cf0b	None	7.0000	12	None	None	None
ce93d206-7144-11e5-ba71-058fbc01cf0b	None	ce8be19a-7144-11e5-ba71-058fbc01cf0b	None	6.7273	11	None	None	None
ce97cd7a-7144-11e5-ba71-058fbc01cf0b	None	ce948926-7144-11e5-ba71-058fbc01cf0b	None	6.7273	11	None	None	None
ce7038dc-7144-11e5-ba71-058fbc01cf0b	None	ce71a230-7144-11e5-ba71-058fbc01cf0b	None	6.6667	18	None	None	None
ce7dc3da-7144-11e5-ba71-058fbc01cf0b	None	ce7cbb2-7144-11e5-ba71-058fbc01cf0b	None	6.3750	16	None	None	None
ce866fda-7144-11e5-ba71-058fbc01cf0b	None	ce80c12a-7144-11e5-ba71-058fbc01cf0b	None	6.3000	10	None	None	None

CODE

```
%%sql SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating) AS AvgRating,
COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type FROM reviews r LEFT JOIN dogs d ON r.dog_guid=d.dog_guid AND
r.user_guid=d.user_guid WHERE d.dog_guid IS NULL GROUP BY r.dog_guid HAVING NumRatings >=10 ORDER BY AvgRating DESC;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 894 rows affected.



In order to make it easier to practice SQL queries with meaningful examples before we learned how to join tables, I added extra columns to the “dogs” table that were not in the original Dognition database. These extra columns included the “total\_tests\_completed” field and multiple inter-test-interval (“iti”) summary fields. **Please do NOT try to use these extra fields in the query exercises below. Since you now know how to join tables, we will practice writing queries as if you only had the data provided in the original Dognition database.**

**Question 3:** How would you use a left join to retrieve a list of all the unique dogs in the dogs table, and retrieve a count of how many tests each one completed? Include the dog\_guids and user\_guids from the dogs and complete\_tests tables in your output. (If you do not limit your query, your output should contain 35050 rows. HINT: use the dog\_guid from the dogs table to group your results.)

In [20]:

```
%%sql
SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name)
FROM dogs d LEFT JOIN complete_tests c
ON d.dog_guid=c.dog_guid
GROUP BY d.dog_guid
LIMIT 10;

* mysql://studentuser:***@localhost/dognitiondb
10 rows affected.
```

Out[20]:

dDogID	cDogID	dUserID	cUserID	COUNT(test_name)
fd27b272-7144-11e5-ba71-058fbc01cf0b	fd27b272-7144-11e5-ba71-058fbc01cf0b	ce134e42-7144-11e5-ba71-058fbc01cf0b	None	21
fd27b5ba-7144-11e5-ba71-058fbc01cf0b	fd27b5ba-7144-11e5-ba71-058fbc01cf0b	ce1353d8-7144-11e5-ba71-058fbc01cf0b	None	20
fd27b6b4-7144-11e5-ba71-058fbc01cf0b	fd27b6b4-7144-11e5-ba71-058fbc01cf0b	ce135ab8-7144-11e5-ba71-058fbc01cf0b	None	2
fd27b79a-7144-11e5-ba71-058fbc01cf0b	fd27b79a-7144-11e5-ba71-058fbc01cf0b	ce13507c-7144-11e5-ba71-058fbc01cf0b	None	11
fd27b86c-7144-11e5-ba71-058fbc01cf0b	fd27b86c-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	None	31
fd27b948-7144-11e5-ba71-058fbc01cf0b	fd27b948-7144-11e5-ba71-058fbc01cf0b	ce13615c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27ba1a-7144-11e5-ba71-058fbc01cf0b	fd27ba1a-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	None	27
fd27baec-7144-11e5-ba71-058fbc01cf0b	None	ce1362ba-7144-11e5-ba71-058fbc01cf0b	None	0
fd27bbbe-7144-11e5-ba71-058fbc01cf0b	fd27bbbe-7144-11e5-ba71-058fbc01cf0b	ce135f2c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27be84-7144-11e5-ba71-058fbc01cf0b	None	ce13697c-7144-11e5-ba71-058fbc01cf0b	None	0

CODE

```
%%sql SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name) FROM dogs d LEFT JOIN complete_tests c ON d.dog_guid=c.dog_guid GROUP BY d.dog_guid;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 35050 rows affected.

Sometimes you can get so focused on writing your join statement that you don't pay close attention to the fields and tables you put in your other clauses, especially when you are joining a lot of tables. Often your query will still run successfully, even if you haven't entered the criteria or grouping clause you intended. The next question will illustrate how easy it is for this to happen.

**Question 4: Repeat the query you ran in Question 3, but intentionally use the dog\_guids from the completed\_tests table to group your results instead of the dog\_guids from the dogs table. (Your output should contain 17987 rows)**

In [23]:

```
%%sql
SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name)
FROM dogs d LEFT JOIN complete_tests c
    ON d.dog_guid=c.dog_guid
GROUP BY c.dog_guid
LIMIT 10;

* mysql://studentuser:***@localhost/dognitiondb
10 rows affected.
```

Out[23]:

dDogID	cDogID	dUserID	cUserID	COUNT(test_name)
fd27baec-7144-11e5-ba71-058fbc01cf0b	None	ce1362ba-7144-11e5-ba71-058fbc01cf0b	None	0
fd27b272-7144-11e5-ba71-058fbc01cf0b	fd27b272-7144-11e5-ba71-058fbc01cf0b	ce134e42-7144-11e5-ba71-058fbc01cf0b	None	21
fd27b5ba-7144-11e5-ba71-058fbc01cf0b	fd27b5ba-7144-11e5-ba71-058fbc01cf0b	ce1353d8-7144-11e5-ba71-058fbc01cf0b	None	20
fd27b6b4-7144-11e5-ba71-058fbc01cf0b	fd27b6b4-7144-11e5-ba71-058fbc01cf0b	ce135ab8-7144-11e5-ba71-058fbc01cf0b	None	2
fd27b79a-7144-11e5-ba71-058fbc01cf0b	fd27b79a-7144-11e5-ba71-058fbc01cf0b	ce13507c-7144-11e5-ba71-058fbc01cf0b	None	11
fd27b86c-7144-11e5-ba71-058fbc01cf0b	fd27b86c-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	None	31
fd27b948-7144-11e5-ba71-058fbc01cf0b	fd27b948-7144-11e5-ba71-058fbc01cf0b	ce13615c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27ba1a-7144-11e5-ba71-058fbc01cf0b	fd27ba1a-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	None	27
fd27bbbe-7144-11e5-ba71-058fbc01cf0b	fd27bbbe-7144-11e5-ba71-058fbc01cf0b	ce135f2c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27c1c2-7144-11e5-ba71-058fbc01cf0b	fd27c1c2-7144-11e5-ba71-058fbc01cf0b	ce136a1c-7144-11e5-ba71-058fbc01cf0b	None	20

CODE

```
%%sql SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name) FROM dogs d
LEFT JOIN complete_tests c ON d.dog_guid=c.dog_guid GROUP BY c.dog_guid;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 17987 rows affected.

This time your query ran successfully, but you retrieved many fewer DogIDs because the GROUP BY clause grouped your results according to the dog\_guids in the completed\_tests table rather than the dog\_guid table. As a result, even though you implemented your join correctly, all of the dog\_guids that were in the dogs table but not in the completed\_tests table got rolled up into one row of your output where completed\_tests.dog\_guid = NULL. This is a good opportunity to remind ourselves about the differences between SELECT/GROUP BY and COUNT DISTINCT.

**Question 5: Write a query using COUNT DISTINCT to determine how many distinct dog\_guids there are in the completed\_tests table.**

In [24]:

```
%%sql
SELECT COUNT(DISTINCT dog_guid)
FROM complete_tests

* mysql://studentuser:***@localhost/dognitiondb
1 rows affected.
```

Out[24]:

COUNT(DISTINCT dog_guid)
17986

CODE

```
%%sql SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name) FROM dogs d
LEFT JOIN complete_tests c ON d.dog_guid=c.dog_guid WHERE c.dog_guid IS NOT NULL GROUP BY c.dog_guid;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 17986 rows affected.



In [8]:

```
%%sql
SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name)
FROM dogs d LEFT JOIN complete_tests c
ON d.dog_guid=c.dog_guid
WHERE d.dog_guid IS NOT NULL
GROUP BY d.dog_guid
LIMIT 20;
```

\* mysql://studentuser:\*\*\*@localhost/dognitiondb  
20 rows affected.

Out[8]:

dDogID	cDogID	dUserID	cUserID	COUNT(test_name)
fd27b272-7144-11e5-ba71-058fbc01cf0b	fd27b272-7144-11e5-ba71-058fbc01cf0b	ce134e42-7144-11e5-ba71-058fbc01cf0b	None	21
fd27b5ba-7144-11e5-ba71-058fbc01cf0b	fd27b5ba-7144-11e5-ba71-058fbc01cf0b	ce1353d8-7144-11e5-ba71-058fbc01cf0b	None	20
fd27b6b4-7144-11e5-ba71-058fbc01cf0b	fd27b6b4-7144-11e5-ba71-058fbc01cf0b	ce135ab8-7144-11e5-ba71-058fbc01cf0b	None	2
fd27b79a-7144-11e5-ba71-058fbc01cf0b	fd27b79a-7144-11e5-ba71-058fbc01cf0b	ce13507c-7144-11e5-ba71-058fbc01cf0b	None	11
fd27b86c-7144-11e5-ba71-058fbc01cf0b	fd27b86c-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	None	31
fd27b948-7144-11e5-ba71-058fbc01cf0b	fd27b948-7144-11e5-ba71-058fbc01cf0b	ce13615c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27ba1a-7144-11e5-ba71-058fbc01cf0b	fd27ba1a-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	None	27
fd27baec-7144-11e5-ba71-058fbc01cf0b	None	ce1362ba-7144-11e5-ba71-058fbc01cf0b	None	0
fd27bbbe-7144-11e5-ba71-058fbc01cf0b	fd27bbbe-7144-11e5-ba71-058fbc01cf0b	ce135f2c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27be84-7144-11e5-ba71-058fbc01cf0b	None	ce13697c-7144-11e5-ba71-058fbc01cf0b	None	0
fd27bf60-7144-11e5-ba71-058fbc01cf0b	None	ce1352ac-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c032-7144-11e5-ba71-058fbc01cf0b	None	ce1352ac-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c0fa-7144-11e5-ba71-058fbc01cf0b	None	ce136a1c-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c1c2-7144-11e5-ba71-058fbc01cf0b	fd27c1c2-7144-11e5-ba71-058fbc01cf0b	ce136a1c-7144-11e5-ba71-058fbc01cf0b	None	20
fd27c294-7144-11e5-ba71-058fbc01cf0b	None	ce1366e8-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c35c-7144-11e5-ba71-058fbc01cf0b	None	ce1366e8-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c424-7144-11e5-ba71-058fbc01cf0b	None	ce1366e8-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c4ec-7144-11e5-ba71-058fbc01cf0b	None	ce1366e8-7144-11e5-ba71-058fbc01cf0b	None	0
fd27c5be-7144-11e5-ba71-058fbc01cf0b	fd27c5be-7144-11e5-ba71-058fbc01cf0b	ce136ac6-7144-11e5-ba71-058fbc01cf0b	None	20
fd27c64a-7144-11e5-ba71-058fbc01cf0b	None	ce136c24-7144-11e5-ba71-058fbc01cf0b	None	0

CODE

```
%%sql SELECT d.dog_guid AS dDogID, c.dog_guid AS cDogID, d.user_guid AS dUserID, c.user_guid AS cUserID, COUNT(test_name) FROM dogs d LEFT JOIN complete_tests c ON d.dog_guid=c.dog_guid WHERE d.dog_guid IS NOT NULL GROUP BY d.dog_guid;
```

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 35050 rows affected.

The result of your COUNT DISTINCT clause should be 17,986 which is one row less than the number of rows you retrieved from your query in Question 4. That's because COUNT DISTINCT does NOT count NULL values, while SELECT/GROUP BY clauses roll up NULL values into one group. If you want to infer the number of distinct entries from the results of a query using joins and GROUP BY clauses, remember to include an "IS NOT NULL" clause to ensure you are not counting NULL values.

These exercises are a good illustration of why it is very helpful to save your queries when you are doing an analysis. Saving your queries allows you and your team members to double-check your work later. As you can see, the concepts behind SQL aren't themselves too tricky, but it is easy to make mistakes, especially when your queries get long and more complicated.

One more situation where joins can cause some confusion is when you have duplicate rows in a table you are joining. If you ignore what we've discussed about set theory and the way databases compute their joins, the behavior databases exhibit when you have duplicate rows in a joined table will seem utterly baffling. With this knowledge, though, the behavior will make perfect sense. Let's walk through what happens.

**Question 6: We want to extract all of the breed information of every dog a user\_guid in the users table owns. If a user\_guid in the users table does not own a dog, we want that information as well. Write a query that would return this information. Include the dog\_guid from the dogs table, and user\_guid from both the users and dogs tables in your output. (HINT: you should get 952557 rows in your output!)**

In [10]:

```
%%sql
SELECT u.user_guid AS uUserID, d.user_guid AS dUserID, d.dog_guid AS dDogID,
d.breed
FROM users u LEFT JOIN dogs d
ON u.user_guid=d.user_guid
LIMIT 10;
```

\* mysql://studentuser:\*\*\*@localhost/dognitiondb  
10 rows affected.

Out[10]:

	uUserID	dUserID	dDogID	breed
	ce134e42-7144-11e5-ba71-058fbc01cf0b	ce134e42-7144-11e5-ba71-058fbc01cf0b	fd27b272-7144-11e5-ba71-058fbc01cf0b	Labrador Retriever
	ce134e42-7144-11e5-ba71-058fbc01cf0b	ce134e42-7144-11e5-ba71-058fbc01cf0b	fd417cac-7144-11e5-ba71-058fbc01cf0b	Mixed
	ce1353d8-7144-11e5-ba71-058fbc01cf0b	ce1353d8-7144-11e5-ba71-058fbc01cf0b	fd27b5ba-7144-11e5-ba71-058fbc01cf0b	Shetland Sheepdog
	ce1353d8-7144-11e5-ba71-058fbc01cf0b	ce1353d8-7144-11e5-ba71-058fbc01cf0b	fd3fb0f2-7144-11e5-ba71-058fbc01cf0b	Shetland Sheepdog
	ce135ab8-7144-11e5-ba71-058fbc01cf0b	ce135ab8-7144-11e5-ba71-058fbc01cf0b	fd27b6b4-7144-11e5-ba71-058fbc01cf0b	Golden Retriever
	ce13507c-7144-11e5-ba71-058fbc01cf0b	ce13507c-7144-11e5-ba71-058fbc01cf0b	fd27b79a-7144-11e5-ba71-058fbc01cf0b	Golden Retriever
	ce135e14-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	fd27b86c-7144-11e5-ba71-058fbc01cf0b	Shih Tzu
	ce135e14-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	fd27ba1a-7144-11e5-ba71-058fbc01cf0b	Shih Tzu
	ce135e14-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	fd27e9a4-7144-11e5-ba71-058fbc01cf0b	Shih Tzu
	ce135e14-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	fd27ed46-7144-11e5-ba71-058fbc01cf0b	Shih Tzu

CODE

```
%%sql SELECT u.user_guid AS uUserID, d.user_guid AS dUserID, d.dog_guid AS dDogID, d.breed FROM dogs d RIGHT JOIN users u ON d.user_guid=u.user_guid;
```

OUTPUT

- \* mysql://studentuser:\*\*\*@localhost/dognitiondb 952557 rows affected.

There are only 35050 distinct dog\_guids in the dogs table. Why is the database outputting almost a million rows? That can't be right. Let's figure out what is going on.

**Question 7: Adapt the query you wrote above so that it counts the number of rows the join will output per user\_id. Sort the results by this count in descending order. Remember that if you include dog\_guid or breed fields in this query, they will be randomly populated by only one of the values associated with a user\_guid (see MySQL Exercise 6; there should be 33,193 rows in your output).**

In [5]:

```
%%sql
SELECT COUNT(DISTINCT dog_guid)
FROM dogs

* mysql://studentuser:***@localhost/dognitiondb
1 rows affected.
```

Out[5]:

```
COUNT(DISTINCT dog_guid)
35050
```

In [7]:

```
%%sql
SELECT u.user_guid AS uUserID, d.user_guid AS dUserID, d.dog_guid AS dDogID, d.breed, count(*) AS numrows
FROM users u LEFT JOIN dogs d
    ON u.user_guid=d.user_guid
GROUP BY u.user_guid
ORDER BY numrows DESC
LIMIT 20;
```

\* mysql://studentuser:\*\*\*@localhost/dognitiondb  
20 rows affected.

Out[7]:

uUserID	dUserID	dDogID	breed	numrows
ce7b75bc-7144-11e5-ba71-058fbc01cf0b	ce7b75bc-7144-11e5-ba71-058fbc01cf0b	fd7bfb52-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	913138
ce225842-7144-11e5-ba71-058fbc01cf0b	ce225842-7144-11e5-ba71-058fbc01cf0b	fd423714-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	442
ce2258a6-7144-11e5-ba71-058fbc01cf0b	ce2258a6-7144-11e5-ba71-058fbc01cf0b	fd40bd62-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	320
ce135e14-7144-11e5-ba71-058fbc01cf0b	ce135e14-7144-11e5-ba71-058fbc01cf0b	fd27b86c-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	130
ce29675e-7144-11e5-ba71-058fbc01cf0b	ce29675e-7144-11e5-ba71-058fbc01cf0b	fd46b014-7144-11e5-ba71-058fbc01cf0b	Labrador Retriever- Mix	110
ce6676d0-7144-11e5-ba71-058fbc01cf0b	ce6676d0-7144-11e5-ba71-058fbc01cf0b	fd68ed6e-7144-11e5-ba71-058fbc01cf0b	Golden Retriever	64
ce7adeea-7144-11e5-ba71-058fbc01cf0b	ce7adeea-7144-11e5-ba71-058fbc01cf0b	fd7b3faa-7144-11e5-ba71-058fbc01cf0b	Labrador Retriever	49
ce47264a-7144-11e5-ba71-058fbc01cf0b	ce47264a-7144-11e5-ba71-058fbc01cf0b	fd401614-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	36
ce8c2d08-7144-11e5-ba71-058fbc01cf0b	ce8c2d08-7144-11e5-ba71-058fbc01cf0b	fdb54786-7144-11e5-ba71-058fbc01cf0b	Basenji	36
ce66713a-7144-11e5-ba71-058fbc01cf0b	ce66713a-7144-11e5-ba71-058fbc01cf0b	fd68e80a-7144-11e5-ba71-058fbc01cf0b	Puggle	30
ce964888-7144-11e5-ba71-058fbc01cf0b	ce964888-7144-11e5-ba71-058fbc01cf0b	fdbbdd4e-7144-11e5-ba71-058fbc01cf0b	Mixed	30
ce26b266-7144-11e5-ba71-058fbc01cf0b	ce26b266-7144-11e5-ba71-058fbc01cf0b	fd436e7c-7144-11e5-ba71-058fbc01cf0b	Pug	25
ce6e67e6-7144-11e5-ba71-058fbc01cf0b	ce6e67e6-7144-11e5-ba71-058fbc01cf0b	fd6cbb42-7144-11e5-ba71-058fbc01cf0b	Australian Shepherd	25
ce728bf0-7144-11e5-ba71-058fbc01cf0b	ce728bf0-7144-11e5-ba71-058fbc01cf0b	fd71a2f6-7144-11e5-ba71-058fbc01cf0b	Labrador Retriever-Golden Retriever Mix	25
ce135766-7144-11e5-ba71-058fbc01cf0b	ce135766-7144-11e5-ba71-058fbc01cf0b	fd3ccf2c-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	24
ce66b9b0-7144-11e5-ba71-058fbc01cf0b	ce66b9b0-7144-11e5-ba71-058fbc01cf0b	fd692aa4-7144-11e5-ba71-058fbc01cf0b	Mixed	24
ce83d2ca-7144-11e5-ba71-058fbc01cf0b	ce83d2ca-7144-11e5-ba71-058fbc01cf0b	fdad96a8-7144-11e5-ba71-058fbc01cf0b	Cockapoo	24
ce7da332-7144-11e5-ba71-058fbc01cf0b	ce7da332-7144-11e5-ba71-058fbc01cf0b	fd80bebc-7144-11e5-ba71-058fbc01cf0b	Mixed	20
ce134492-7144-11e5-ba71-058fbc01cf0b	ce134492-7144-11e5-ba71-058fbc01cf0b	fd40e206-7144-11e5-ba71-058fbc01cf0b	Shih Tzu	18
ce9a381c-7144-11e5-ba71-058fbc01cf0b	ce9a381c-7144-11e5-ba71-058fbc01cf0b	fdc12bbe-7144-11e5-ba71-058fbc01cf0b	Chihuahua	18

CODE

%%sql SELECT u.user\_guid AS uUserID, d.user\_guid AS dUserID, d.dog\_guid AS dDogID, d.breed, count(\*) AS numrows FROM users u LEFT JOIN dogs d ON u.user\_guid=d.user\_guid GROUP BY u.user\_guid ORDER BY numrows DESC;

OUTPUT

- mysql://studentuser:\*\*\*@localhost/dognitiondb 33193 rows affected.

This query told us that user 'ce7b75bc-7144-11e5-ba71-058fbc01cf0b' would be associated with 913,138 rows in the output of the outer join we designed! Once again, why? We are going to work with the second user\_guid in the output you just generated, 'ce225842-7144-11e5-ba71-058fbc01cf0b', because it would be associated with 442 output rows, and 442 rows are much easier to work with than 913,138.

**Question 8: How many rows in the *users* table are associated with user\_guid 'ce225842-7144-11e5-ba71-058fbc01cf0b'?**

In [8]:

```
%%sql
SELECT COUNT(user_guid)
FROM users
WHERE user_guid='ce225842-7144-11e5-ba71-058fbc01cf0b'
```

\* mysql://studentuser:\*\*\*@localhost/dognitiondb  
1 rows affected.

Out[8]:

COUNT(user\_guid)  
17

There are 17 entries associated with that user\_guid in the users table. If you examine all the columns in the entries, you will see that the rows are exact duplicates of each other. That's unfortunate, but also something that can happen in real life data sets, especially those from new companies or governmental agencies.

Ok, now...

Question 9: Examine all the rows in the *dogs* table that are associated with user\_guid 'ce225842-7144-11e5-ba71-058fbc01cf0b'?

In [9]:

```
%%sql
SELECT *
FROM dogs
WHERE user_guid='ce225842-7144-11e5-ba71-058fbc01cf0b'
```

\* mysql://studentuser:\*\*\*@localhost/dognitiondb  
26 rows affected.

Out[9]:

gender	birthday	breed	weight	dog_fixed	dna_tested	created_at	updated_at	dimension	exclude	breed_type	breed_group	dimension
male	2013	Shih Tzu	190	1	1	2013-03-21 19:30:06	2015-05-13 18:01:39	maverick	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-03-30 19:41:08	2015-01-06 16:26:41	stargazer	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-04-08 15:44:08	2015-09-21 18:05:51	einstein	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-05-03 19:56:17	2013-07-25 19:42:39	None	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-05-16 16:03:54	2014-10-20 18:10:48	None	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-05-17 16:03:57	2014-11-05 16:18:29	stargazer	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-05-23 12:29:28	2014-10-06 19:29:11	None	1	Pure Breed	Toy	71
female	2013	Shih Tzu	190	1	1	2013-07-18 19:16:30	2014-05-28 14:11:10	None	1	Pure Breed	Toy	71
male	2013	Shih Tzu	0	1	1	2013-07-30 19:39:37	2013-07-30 19:39:37	None	None	Pure Breed	Toy	71
male	2013	Shih Tzu	190	1	1	2013-08-05 19:50:55	2014-07-20 15:06:18	None	1	Pure Breed	Toy	71
female	2013	Shih Tzu	190	1	1	2013-07-18 17:44:47	2013-10-31 18:32:39	None	1	Pure Breed	Toy	71
male	2013	Shih Tzu	190	0	0	2013-09-24	2014-05-28	None	1	Pure Breed	Toy	71



Set 1		Set 2	
EmployeeID	Name	DepartID	Department
AA	Lucy	1	Clothing
BB	Joe	2	Electronics
CC	Luke		

×

Cartesian (Cross) Product			
AA	Lucy	1	Clothing
AA	Lucy	2	Electronics
BB	Joe	1	Clothing
BB	Joe	2	Electronics
CC	Luke	1	Clothing
CC	Luke	2	Electronics

When databases join tables, they output the result of every pair of entries that meet certain criteria in the linking column of one table with the linking column of another table. Our join statement imposed the criteria that the output should only include pairs whose user\_guids matched in the two linking columns. However, since there were multiple rows that had the same user\_guid in the users table, *each one of these rows got paired up with each row in the dogs table that had the same user\_guid*. The result was 442 rows, because 17 (instances of the user\_guid in the users table) x 26 (instances of the user\_guid in the dogs table) = 442.

Having seen this, perhaps you can now appreciate why some database experts emphasize terminology that differentiates between set *theory* and real database *implementation*. Database operations, like those that join tables, are based on set theory that assumes there are no duplicate rows in your tables. In real life, duplicate rows get entered all the time, and it can cause you to misinterpret your queries if you do not understand the consequences. If you've been impacted enough times by the differences between real and theoretical databases, it makes sense that it would be important to you to use language that clearly distinguishes between theory and real life.

The important things I want you to remember from this example of joins with duplicates are that duplicate rows and table relationships that have table-to-table mappings of greater than 1 have multiplicative effects on your query results, due to the way relational databases combine tables. If you write queries that aggregate over a lot of joined tables, it can be very difficult to catch issues that output results you don't intend, because the aggregated results will hide clues from you. To prevent this from happening, I recommend you adopt the following practices:

- Avoid making assumptions about your data or your analyses. For example, rather than assume that all the values in a column are unique just because some documentation says they should be, check for yourself!
- Always look at example outputs of your queries before you strongly interpret aggregate calculations. Take extra care to do this when your queries require joins.
- When your queries require multiple layers of functions or joins, examine the output of each layer or join first before you combine them all together.
- Adopt a healthy skepticism of all your data and results. If you see something you don't expect, make sure you explore it before interpreting it strongly or incorporating it into other analyses.

One more type of join to mention that I discussed in the joins videos is a full outer join. Full outer joins include all of the rows in both tables in an ON clause, regardless of whether there is a value that links the row of one table with a row in the other table. As with left or right joins, whenever a value in a row does not have a matching value in the joined table, NULLs will be entered for all values in the joined table.

Outer joins are used very rarely. The most practical application is if you want to export all of your raw data to another program for visualization or analysis. The syntax for outer joins is the same as for inner joins, but you replace the word "inner" with "full outer":

```
SELECT r.dog_guid AS rDogID, d.dog_guid AS dDogID, r.user_guid AS rUserID, d.user_guid AS dUserID, AVG(r.rating) AS AvgRating, COUNT(r.rating) AS NumRatings, d.breed, d.breed_group, d.breed_type
FROM reviews r FULL OUTER JOIN dogs d
ON r.dog_guid=d.dog_guid AND r.user_guid=d.user_guid
WHERE r.dog_guid IS NOT NULL
GROUP BY r.dog_guid
ORDER BY AvgRating DESC;
```

**\*\*HOWEVER! MySQL does not support full outer joins.\*\***

If you wanted to imitate a full outer join in MySQL, you could follow one of the methods described at this website:

<http://www.xaprb.com/blog/2006/05/26/how-to-write-full-outer-join-in-mysql/> (<http://www.xaprb.com/blog/2006/05/26/how-to-write-full-outer-join-in-mysql/>)

## Practice outer joining your own tables!¶

**Question 10:** How would you write a query that used a *left* join to return the number of distinct user\_guids that were in the users table, but not the dogs table (your query should return a value of 2226)?

In [5]:

```
%%sql
SELECT COUNT(DISTINCT u.user_guid)
FROM users u LEFT JOIN dogs d
ON u.user_guid=d.user_guid
WHERE d.user_guid IS NULL;

* mysql://studentuser:***@localhost/dognitiondb
1 rows affected.
```

Out[5]:

```
COUNT(DISTINCT u.user_guid)
2226
```

**Question 11:** How would you write a query that used a *right* join to return the number of distinct user\_guids that were in the users table, but not the dogs table (your query should return a value of 2226)?

In [6]:

```
%%sql
SELECT COUNT(DISTINCT u.user_guid)
FROM dogs d RIGHT JOIN users u
ON d.user_guid=u.user_guid
WHERE d.user_guid IS NULL;

* mysql://studentuser:***@localhost/dognitiondb
1 rows affected.
```

Out[6]:

```
COUNT(DISTINCT u.user_guid)
2226
```

**Question 12:** Use a left join to create a list of all the unique dog\_guids that are contained in the site\_activities table, but not the dogs table, and how many times each one is entered. Note that there are a lot of NULL values in the dog\_guid of the site\_activities table, so you will want to exclude them from your list. (Hint: if you exclude null values, the results you get will have two rows with words in their site\_activities dog\_guid fields instead of real guids, due to mistaken entries)

In [7]:

```
%%sql
SELECT s.dog_guid AS SA_dogs_not_present_in_dogs_table, COUNT(*) AS
NumEntries
FROM site_activities s LEFT JOIN dogs d
ON s.dog_guid=d.dog_guid
WHERE d.dog_guid IS NULL AND s.dog_guid IS NOT NULL
GROUP BY SA_dogs_not_present_in_dogs_table;

* mysql://studentuser:***@localhost/dognitiondb
2 rows affected.
```

Out[7]:

SA_dogs_not_present_in_dogs_table	NumEntries
Membership	5587
PortalContent	12

**Practice any other outer joins you are interested in here!**

In [ ]: